


1|C flowers B50: only S-REND floor fail (-0.15); diffuse speckle, FT can't repaint



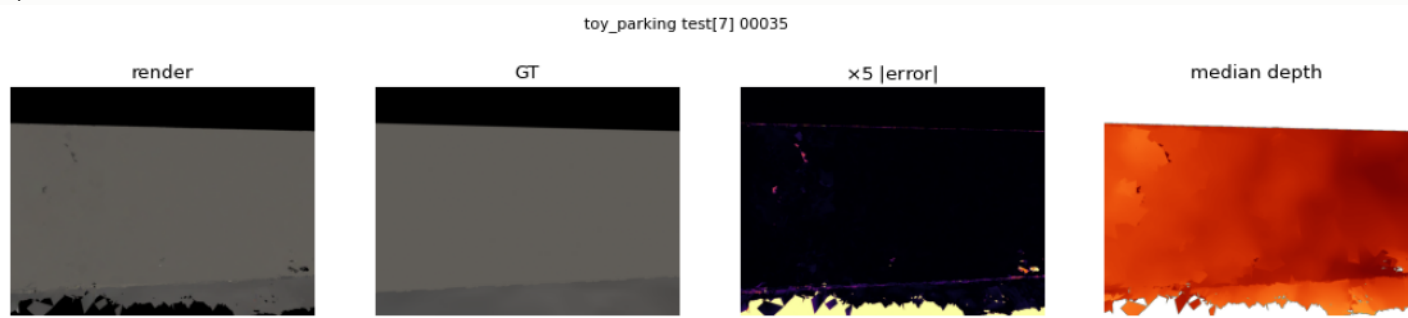
2|C town06 B50 (-0.45): least over-parameterized town; far-field damage



3|A toy B50 residual (-0.52): train-coverage gap erodes wall/ground junction



4|A E2v3 floater deletion: -8.07 dB hole at view 00035 (selection effect)



5|B default FT after prune: -2.9 dB global washout (position drift)



6|B gradient-routed geometry losses: silhouette halos; guard blown



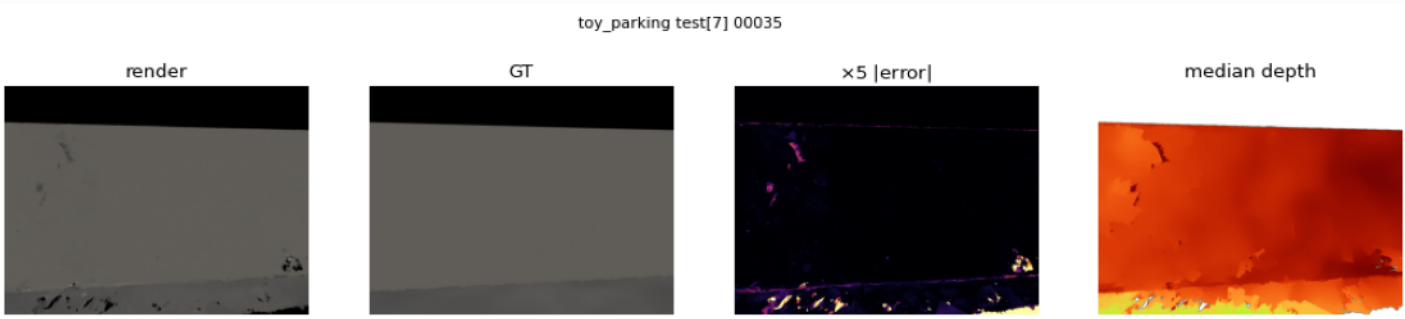
7|D SH-distill on 72-view toy: LPIPS 0.11->0.21 (view-manifold overfit)



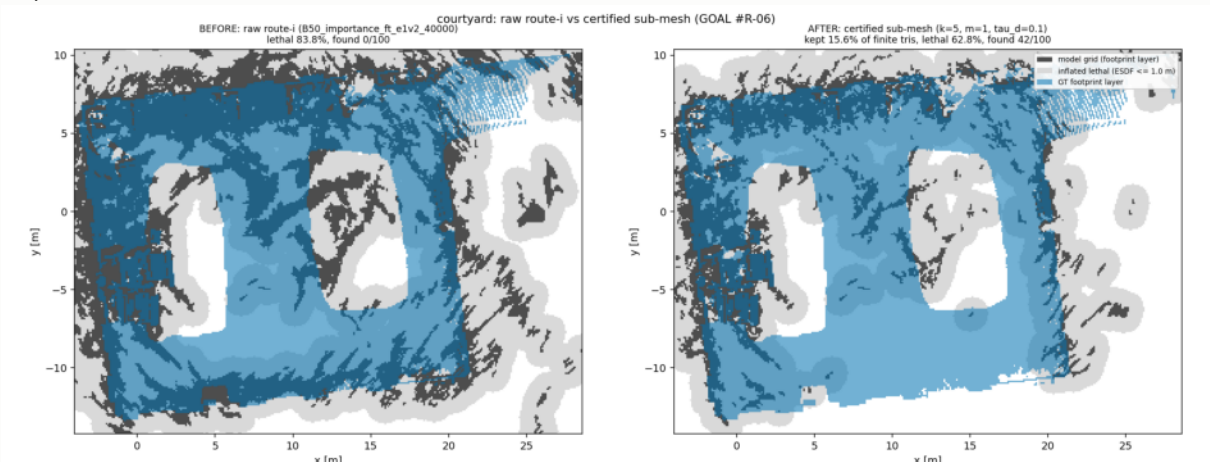
8|C courtyard random-B25 (-3.97): starvation without evidence ranking



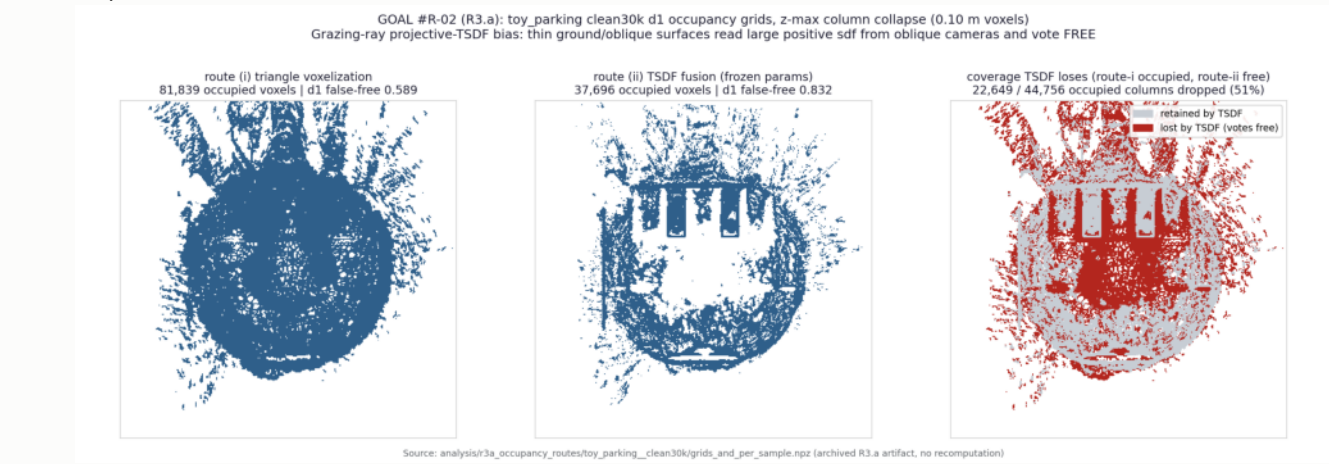
9|A E2R fading on toy: pinholes on low-train-support content; guard blown



10|A/E R3.b certification sheds real walls with junk (coll. 16.7/100)



11|E R3.a TSDF grazing-ray bias: near-surface voxels vote FREE



12|E garden g3 fragmentation 1941->5724 components at iso-PSNR



13|C kitchen B12.5 (-2.37): aggressive-budget starvation, indoor worst



E9-FAIL taxonomy — 5 mechanism families: A train-coverage selection effects · B optimizer drift on converged ckpts · C budget starvation / headroom exhaustion · D view-conditioning capacity mismatch · E baseline-geometry unreliability & consumption-route bias.

13 cases, 24 panels; one-paragraph diagnoses, evidence pointers and 'what it bounds' lines in analysis/e9_failure_taxonomy/TAXONOMY.md (LEDGER GOAL#R-09).